

```
In [2]: #Installing the Libraries
import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: #Importing insurance.csv dataset
insurance_dataset=pd.read_csv("insurance.csv")
```

```
In [4]: insurance_dataset
```

```
Out[4]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

```
In [6]: insurance_dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         1338 non-null   int64
 1   sex         1338 non-null   object
 2   bmi         1338 non-null   float64
 3   children    1338 non-null   int64
 4   smoker      1338 non-null   object
 5   region      1338 non-null   object
 6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

In [7]: `insurance_dataset.describe()`

Out[7]:

	age	bmi	children	charges
<b>count</b>	1338.000000	1338.000000	1338.000000	1338.000000
<b>mean</b>	39.207025	30.663397	1.094918	13270.422265
<b>std</b>	14.049960	6.098187	1.205493	12110.011237
<b>min</b>	18.000000	15.960000	0.000000	1121.873900
<b>25%</b>	27.000000	26.296250	0.000000	4740.287150
<b>50%</b>	39.000000	30.400000	1.000000	9382.033000
<b>75%</b>	51.000000	34.693750	2.000000	16639.912515
<b>max</b>	64.000000	53.130000	5.000000	63770.428010

In [8]: `insurance_dataset.isnull().sum()`

```
Out[8]: age      0  
sex      0  
bmi      0  
children 0  
smoker   0  
region   0  
charges  0  
dtype: int64
```

```
In [9]: insurance_dataset['sex'].value_counts()
```

```
Out[9]: male      676  
female    662  
Name: sex, dtype: int64
```

```
In [10]: insurance_dataset['age'].value_counts()
```

```
Out[10]:
```

18	69
19	68
50	29
51	29
47	29
46	29
45	29
20	29
48	29
52	29
22	28
49	28
54	28
53	28
21	28
26	28
24	28
25	28
28	28
27	28
23	28
43	27
29	27
30	27
41	27
42	27
44	27
31	27
40	27
32	26
33	26
56	26
34	26
55	26
57	26
37	25
59	25
58	25
36	25
38	25
35	25
39	25
61	23
60	23

```
63    23
62    23
64    22
Name: age, dtype: int64
```

```
In [11]: insurance_dataset['children'].value_counts()
```

```
Out[11]: 0    574
         1    324
         2    240
         3    157
         4     25
         5     18
Name: children, dtype: int64
```

```
In [12]: insurance_dataset['smoker'].value_counts()
```

```
Out[12]: no    1064
         yes     274
Name: smoker, dtype: int64
```

```
In [13]: insurance_dataset['region'].value_counts()
```

```
Out[13]: southeast    364
         southwest    325
         northwest    325
         northeast    324
Name: region, dtype: int64
```

```
In [14]: insurance_dataset.columns
```

```
Out[14]: Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges'], dtype='object')
```


```
In [8]: #Data Analysis

#Age vs Charges
plt.figure(figsize = (12, 8))
x=insurance_dataset['age']
y=insurance_dataset['charges']
sns.barplot(x,y)
```

```
plt.title('AGE vs CHARGES')
plt.show()
```

C:\Python\Python38\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```



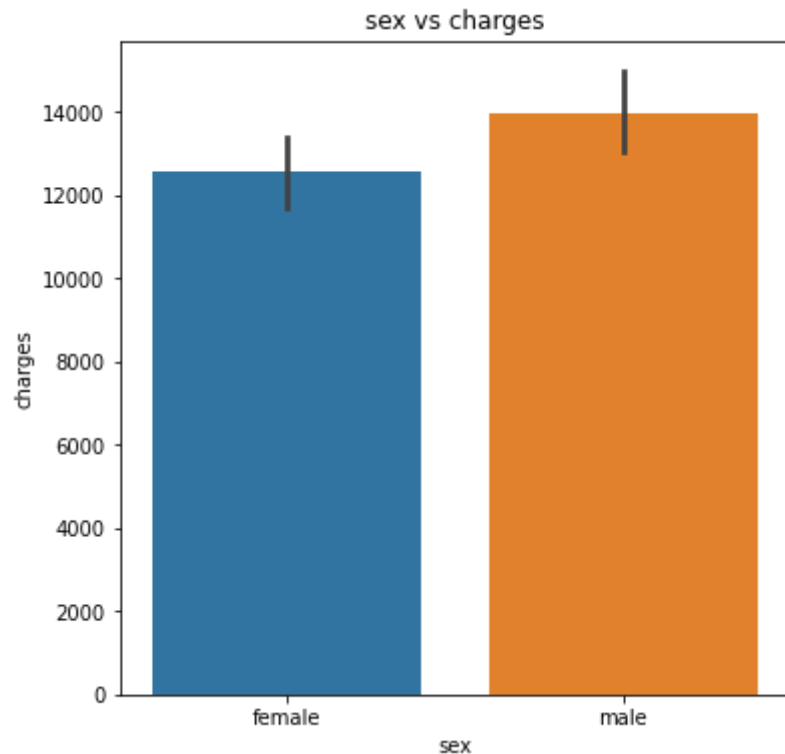
```
In [9]: # sex vs charges
# males have slightly greater insurance charges than females in general

plt.figure(figsize = (6, 6))
```

```
x=insurance_dataset['sex']  
y=insurance_dataset['charges']  
  
sns.barplot(x,y)  
  
plt.title('sex vs charges')  
plt.show()
```

C:\Python\Python38\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(



In [10]:


```
# children vs charges  
# no. of childrens of a person has a very interesting dependency on insurance costs  
  
plt.figure(figsize = (12, 8))  
x=insurance_dataset['children']  
y=insurance_dataset['charges']
```

```
sns.barplot(x,y)

plt.title('children vs charges')
plt.show()
```

C:\Python\Python38\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```



```
In [11]: # region vs charges
```



*# From the graph we can see that the region actually does not play any role in determining the insurance charges*


```
plt.figure(figsize = (12, 8))
x=insurance_dataset['children']
y=insurance_dataset['charges']

sns.barplot(x,y)

plt.title('region vs charges')
plt.show()
```

C:\Python\Python38\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```



In [12]:


```
# smoker vs charges
# from the graph below, it is visible that smokers have more insurance charges than the non smokers

plt.figure(figsize = (6, 6))
x=insurance_dataset['smoker']
y=insurance_dataset['charges']
sns.barplot(x,y)

plt.title('smoker vs charges')
plt.show()
```

C:\Python\Python38\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```



In [14]:

```
# plotting the correlation plot for the dataset
```

```
f, ax = plt.subplots(figsize = (10, 10))
```

```
corr = insurance_dataset.corr()
```

```
sns.heatmap(corr, mask = np.zeros_like(corr, dtype = np.bool),
```


```
            cmap = sns.diverging_palette(50, 10, as_cmap = True), square = True, ax = ax)
```

C:\Users\TOSHIBA\AppData\Local\Temp\ipykernel\_9496\3944541768.py:6: DeprecationWarning: `np.bool` is a deprecated alias for the builtin `bool`. To silence this warning, use `bool` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.bool\_` here.

Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
sns.heatmap(corr, mask = np.zeros_like(corr, dtype = np.bool),
```

Out[14]: &lt;AxesSubplot:&gt;



```
In [15]: # removing unnecassary columns from the dataset  
  
insurance_dataset = insurance_dataset.drop('region', axis = 1)  
  
print(insurance_dataset.shape)
```

```
insurance_dataset.columns
```

```
(1338, 6)
```

```
Out[15]: Index(['age', 'sex', 'bmi', 'children', 'smoker', 'charges'], dtype='object')
```

```
In [17]:
```

```
# Label encoding for sex and smoker

# importing Label encoder
from sklearn.preprocessing import LabelEncoder

# creating a Label encoder
le = LabelEncoder()

# Label encoding for sex
# 0 for females and 1 for males
insurance_dataset['sex'] = le.fit_transform(insurance_dataset['sex'])

# Label encoding for smoker
# 0 for smokers and 1 for non smokers
insurance_dataset['smoker'] = le.fit_transform(insurance_dataset['smoker'])
```

```
In [18]:
```

```
insurance_dataset['sex'].value_counts()
```

```
Out[18]:
```

```
1    676
0    662
Name: sex, dtype: int64
```

```
In [19]:
```

```
insurance_dataset['smoker'].value_counts()
```

```
Out[19]:
```

```
0    1064
1     274
Name: smoker, dtype: int64
```

```
In [22]:
```

```
# splitting the dependent and independent variable

x = insurance_dataset.iloc[:,5]
y = insurance_dataset.iloc[:,5]
```

```
print(x.shape)
print(y.shape)
```

```
(1338, 5)
(1338,)
```

In [24]:

```
# splitting the dataset into training and testing sets

from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2, random_state = 30)

print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(1070, 5)
(268, 5)
(1070,)
(268,)
```

In [25]:

```
#Modelling

#Linear Regression

from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score

# creating the model
model = LinearRegression()

# feeding the training data to the model
model.fit(x_train, y_train)

# predicting the test set results
y_pred = model.predict(x_test)

# calculating the mean squared error
mse = np.mean((y_test - y_pred)**2, axis = None)
```

```
print("MSE :", mse)

# Calculating the root mean squared error
rmse = np.sqrt(mse)
print("RMSE :", rmse)

# Calculating the r2 score
r2 = r2_score(y_test, y_pred)
print("r2 score :", r2)
```

```
MSE : 37057975.783359274
RMSE : 6087.526244983202
r2 score : 0.759758528014896
```

In [26]:

```
y_test
```

Out[26]:

```
338    41919.09700
620     3659.34600
965     4746.34400
128    32734.18630
329     9144.56500
...
580    12913.99240
786    12741.16745
321    24671.66334
903     8125.78450
613     6753.03800
Name: charges, Length: 268, dtype: float64
```

In [27]:

```
y_pred
```

```
Out[27]: array([35206.95988424, 5913.25371451, 5902.11903681, 26652.44561859,
13037.68322943, 6976.63677636, 4714.42087021, 9784.76464047,
33699.76519397, 11386.17430622, 8038.6451187 , 10230.0795605 ,
35207.19003322, 11538.33660241, 1379.9011798 , 36009.29793685,
11972.55413556, 9251.68230172, 29236.87908302, 11307.31833466,
1438.78621227, 8722.85417 , 8119.68124443, 11284.22664405,
33324.77761469, 39447.6034078 , 14940.85243651, 3423.64409705,
11041.14918239, 14793.12711084, 2187.2244934 , 31013.67789827,
1465.9188615 , 15025.76344263, 12731.54461399, 8998.82427631,
2405.26396426, 6992.56675683, 724.24627861, 8735.58652147,
6939.69247432, 34200.98347795, 2536.40883799, 14066.35901945,
10715.10391391, 30446.39967614, 5420.74682483, 3134.91154318,
2864.74903307, 5583.45182778, 11557.18128072, 12987.4115041 ,
12420.44483284, 1407.03382903, 11451.7642406 , 7496.42488881,
31522.42980901, 3573.16654523, 12963.45669772, 5015.13453229,
15235.0742298 , 11213.15074468, 11460.11010792, 7949.0636769 ,
-466.65361082, 8241.77639306, 10298.61718875, 1649.75256961,
3006.35193724, 15048.85142823, 10915.66933214, 2398.74902057,
294.41833049, 9764.87488799, 16276.17413659, -668.29539324,
12728.61382067, 9823.57715698, 31074.04522796, 4210.00456841,
1090.07906619, 11224.20445106, 9559.40169037, 31013.16340893,
1220.13786991, 16889.04045283, 9582.0497683 , 17464.36679108,
9229.03051879, 11296.93510119, 14128.87209807, 13800.56919382,
26876.29247035, 11030.42741716, 37112.86022059, 5720.08615622,
10022.76185487, 1486.82090741, 7150.65202577, 35620.29227205,
34904.1721693 , 8618.65198729, 9969.88757236, 8334.46168583,
33327.58703064, 26229.34215756, 7319.09663385, 36219.40755375,
33958.86041332, 2137.49054761, 10466.7094654 , 6852.63261838,
842.56463401, 28219.28688021, 5983.40570178, 4206.47460942,
12757.02910307, 10159.83689608, 10056.0169256 , 15214.6260922 ,
4465.70189058, 12626.15166915, 36508.91544617, 14342.5798803 ,
29835.75740984, 10946.16978243, 13980.73756495, 8818.76779201,
31340.01781708, 12291.45890359, 11399.9694375 , 5599.01628127,
34078.41418993, 27318.72049284, 12377.22953576, 2996.6049593 ,
40826.47213082, 25519.90286143, 11933.17025369, 6692.04326826,
12707.36685336, 10447.66243841, 3634.87418919, 37619.58464454,
8000.79440925, 10914.13964943, 6571.8635043 , 3553.82508322,
9849.83638049, 32982.15559928, 16134.58503209, 6630.67705599,
11135.77707034, 6363.86174138, 9499.13244785, 8423.33578014,
5115.63408719, 12824.4769563 , 17935.15610553, 5466.63615986,
10799.19219598, 31715.39770325, 14406.3790818 , 10270.82457752,
15338.84681479, 15542.22164896, 34742.01894654, 8224.02889979,
7206.25393349, 2093.20237606, 8533.08450889, 11961.73449846,
6079.68485076, 4885.9790351 , 33494.84028682, 7413.50756836,
```



```

3428.18264445, 10443.50649208, 14845.70695833, 30251.01859315,
10710.05746776, 14359.86397471, 1527.75878257, 581.96672718,
39139.13606464, 28076.22867404, 27336.55309301, 7125.86478951,
31653.19966523, 14799.8684985 , -317.13116264, 9585.74559021,
14102.05448943, 12017.47178418, 1983.25398327, 12623.24438134,
15827.42773357, 31508.63838275, 30456.37680305, 5752.4818162 ,
10147.97426532, 3088.80934855, 7727.54533755, 7179.26675692,
28057.49298264, 1339.18047336, 5099.13852687, 37067.88138681,
35597.4849218 , 37339.03928978, 6680.41449149, 7540.69986481,
39221.15007856, 6749.76704493, 5807.07884052, 38666.85806332,
5892.52123654, 2964.42646805, 16984.90647405, 28616.60090621,
10393.99281591, 38416.15549267, 12139.72620514, 1824.50875233,
3135.80186945, 9975.92221665, 9433.91153017, 7503.13618036,
7527.15178307, 3233.69886717, 9573.98845664, 11188.30953233,
10215.98050366, 13241.2919176 , 9970.93695567, 31930.81422667,
30300.50527411, 11524.54888115, 7471.93208759, 28339.97432762,
16622.83400972, 14994.93497148, 9195.53210353, 6020.93188011,
16319.75185976, 11243.89475479, 1685.52901091, 4565.39252111,
5548.69406955, 15496.09884877, 4821.55716243, 10765.44302618,
13989.92964748, 37230.63028553, 25491.62934668, 4505.607283 ,
11958.15693869, 11360.36158099, 11147.99432833, 27066.27189017,
9734.1104876 , 15590.67935122, 5903.98415038, 11803.17203101,
15255.93898488, 5938.44747358, 12282.81489623, 4325.27016334]])

```

In [28]:

```

#Support Vector Machine

from sklearn.svm import SVR

# creating the model
model = SVR()

# feeding the training data to the model
model.fit(x_train, y_train)

# predicting the test set results
y_pred = model.predict(x_test)

# calculating the mean squared error
mse = np.mean((y_test - y_pred)**2, axis = None)
print("MSE :", mse)

# Calculating the root mean squared error
rmse = np.sqrt(mse)

```

```
print("RMSE :", rmse)

# Calculating the r2 score
r2 = r2_score(y_test, y_pred)
print("r2 score :", r2)
```

```
MSE : 175634760.8944562
RMSE : 13252.726545675656
r2 score : -0.13861463280418374
```

In [29]:

```
#Random Forest

from sklearn.ensemble import RandomForestRegressor

# creating the model
model = RandomForestRegressor(n_estimators = 40, max_depth = 4, n_jobs = -1)

# feeding the training data to the model
model.fit(x_train, y_train)

# predicting the test set results
y_pred = model.predict(x_test)

# calculating the mean squared error
mse = np.mean((y_test - y_pred)**2, axis = None)
print("MSE :", mse)

# Calculating the root mean squared error
rmse = np.sqrt(mse)
print("RMSE :", rmse)

# Calculating the r2 score
r2 = r2_score(y_test, y_pred)
print("r2 score :", r2)
```

```
MSE : 23282926.865853906
RMSE : 4825.23852942566
r2 score : 0.8490601684486496
```

In [30]:

```
print(y_pred)
```

[45023.8249976 5639.67721012 6384.16753679 16871.65933407  
13028.5036422 5836.29716648 2551.33841792 10583.17884479  
45842.30371458 13884.64904764 7027.60150419 10529.80632979  
45369.70039144 13024.85448321 2551.33841792 24736.10876769  
10498.11468957 10294.73817716 17724.37010433 10573.49089461  
2505.86411698 6560.96787865 10406.79019751 9666.72650457  
25750.57493142 45592.57965972 13990.73313932 5781.70786169  
7261.46255252 13808.74403919 5847.33933626 38360.81721999  
4763.01071051 13567.3010879 14183.21630378 6560.96787865  
2742.57889882 7027.60150419 2505.86411698 6079.18113514  
6179.6715885 23924.40784553 4759.5571758 13211.36538649  
7197.20446162 39255.30714911 6179.6715885 2551.33841792  
5282.95898671 6375.53819724 8264.64012551 10445.91190019  
13925.09462838 2570.67545494 7334.51026652 9189.551811  
26503.63393383 5217.67915095 14183.21630378 6263.05825556  
13211.36538649 6718.31065152 10912.22503258 5557.47129699  
2505.86411698 5639.67721012 10692.81692405 5847.33933626  
6608.40410149 13211.36538649 14471.99020055 2677.76756087  
2961.8105536 10602.94195693 15616.21088009 3325.744939  
10157.47123256 10692.81692405 24315.78164489 4931.46061968  
2723.24186181 10293.81236002 6275.74363 18839.3202548  
2723.24186181 14678.67561072 6879.2414848 14678.67561072  
10445.0349841 7106.03531419 14394.10165927 13621.09215416  
36737.28421625 8220.36295948 26652.88211913 4915.68753843  
10615.9196883 6711.06516379 6722.84370049 45858.44006512  
45023.8249976 13200.46339318 7367.01372047 5803.72004579  
38523.85418439 17491.22707563 6793.30136319 45023.8249976  
25750.57493142 4716.64918416 13396.70903154 7282.80151318  
6236.86278904 18688.22316658 5039.25978079 6138.288779  
13211.36538649 6783.58085694 9936.14526458 10511.82855112  
2961.8105536 7766.90298964 25945.40353823 14048.38755868  
18839.3202548 14662.21926771 14584.33072643 9313.78622073  
25482.66705192 13401.59445365 13824.00443605 4915.68753843  
40623.579942 36464.19853444 13401.59445365 2570.67545494  
45692.20781011 18449.90486539 14295.74622458 6270.01579517  
13215.08355037 10660.07608054 5863.91377483 45592.57965972  
7005.91979204 13633.77536889 5675.10471918 6600.52810956  
9426.22786156 25133.70261607 13401.59445365 10330.16397071  
6372.29887344 9470.13454269 7100.64921818 6877.90195749  
2723.24186181 13401.59445365 13804.36742526 2723.24186181  
10610.81606687 37852.2824917 10384.66527525 10716.90613401  
10649.70140286 14768.02962255 41600.03644142 6722.84370049  
5717.43225794 4983.7937405 7213.26760507 13401.59445365  
7230.15635675 5639.67721012 38250.44883616 10256.49339057

```

2505.86411698 5412.44651803 14584.33072643 24315.78164489
10778.19367322 13401.59445365 2742.57889882 2742.57889882
45592.57965972 39255.30714911 16871.65933407 6318.37609115
38092.96630134 13066.21042425 2551.33841792 5558.76888791
14584.33072643 13759.14813112 2723.24186181 13401.59445365
15806.43994725 24181.54845896 37883.3934059 5813.91754985
7164.90730909 2505.86411698 7100.64921818 6026.5359787
38435.53379463 6481.81394692 4888.55262804 45592.57965972
45616.41427775 42129.65362918 6270.01579517 6739.45614335
45692.20781011 10330.16397071 5610.48131784 45692.20781011
5039.25978079 2677.76756087 13621.09215416 18983.78150065
6952.28919879 45592.57965972 9546.85719044 2570.67545494
5263.49634549 8297.85279644 7027.60150419 4993.44258625
7027.60150419 2723.24186181 7370.91591337 13462.58740549
10716.90613401 13401.59445365 10635.41179149 36969.04285858
39255.30714911 13401.59445365 6667.24362834 36837.26159522
14488.44654356 15620.07423318 13166.47477502 2742.57889882
15383.00789583 14662.21926771 2723.24186181 4636.01037091
5039.25978079 13211.36538649 6306.67192811 10667.95019048
10713.15578601 25945.40353823 20657.49595041 6059.96410782
10634.73751955 13024.85448321 13812.75410249 21530.45804271
7027.60150419 14584.33072643 5117.01482861 14490.51770446
14277.56118807 6263.05825556 10307.7680395 6963.69270806]

```

In [31]:

```
print(y_test)
```

```

338    41919.09700
620     3659.34600
965     4746.34400
128    32734.18630
329     9144.56500
...
580    12913.99240
786    12741.16745
321    24671.66334
903     8125.78450
613     6753.03800

```

```
Name: charges, Length: 268, dtype: float64
```

In [32]:

```
#Decision Forest
```

```
from sklearn.tree import DecisionTreeRegressor
```

```
# creating the model
model = DecisionTreeRegressor()

# feeding the training data to the model
model.fit(x_train, y_train)

# predicting the test set results
y_pred = model.predict(x_test)

# calculating the mean squared error
mse = np.mean((y_test - y_pred)**2, axis = None)
print("MSE :", mse)

# Calculating the root mean squared error
rmse = np.sqrt(mse)
print("RMSE :", rmse)

# Calculating the r2 score
r2 = r2_score(y_test, y_pred)
print("r2 score :", r2)
```

MSE : 44043931.24033023

RMSE : 6636.560196391669

r2 score : 0.7144695939399076

In [ ]: